

# softmax

September 28, 2018

## 1 Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
In [1]: import random
import numpy as np
from cs682.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

from __future__ import print_function

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

In [2]: def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000, num_dev=500):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the linear classifier. These are the same steps as we used for the
    SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
```

```

cifar10_dir = 'cs682/datasets/cifar-10-batches-py'

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# subsample the data
mask = list(range(num_training, num_training + num_validation))
X_val = X_train[mask]
y_val = y_train[mask]
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# Normalize the data: subtract the mean image
mean_image = np.mean(X_train, axis = 0)
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

# add bias dimension and transform into columns
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev

# Cleaning up variables to prevent loading data multiple times (which may cause memory
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

```

```

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
print('dev data shape: ', X_dev.shape)
print('dev labels shape: ', y_dev.shape)

```

```

Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)

```

## 1.1 Softmax Classifier

Your code for this section will all be written inside `cs682/classifiers/softmax.py`.

```

In [3]: # First implement the naive softmax loss function with nested loops.
# Open the file cs682/classifiers/softmax.py and implement the
# softmax_loss_naive function.

from cs682.classifiers.softmax import softmax_loss_naive
import time

# Generate a random softmax weight matrix and use it to compute the loss.
W = np.random.randn(3073, 10) * 0.0001
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As a rough sanity check, our loss should be something close to -log(0.1).
print('loss: %f' % loss)
print('sanity check: %f' % (-np.log(0.1)))

```

```

loss: 2.355978
sanity check: 2.302585

```

## 1.2 Inline Question 1:

Why do we expect our loss to be close to  $-\log(0.1)$ ? Explain briefly.\*\*

**Your answer:** The probability of each class 0.1, because there are 10 classes. Hence on an average, we can expect the loss to be  $-\log(0.1)$

```
In [4]: # Complete the implementation of softmax_loss_naive and implement a (naive)
# version of the gradient that uses nested loops.
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As we did for the SVM, use numeric gradient checking as a debugging tool.
# The numeric gradient should be close to the analytic gradient.
from cs682.gradient_check import grad_check_sparse
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)

# similar to SVM case, do another gradient check with regularization
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

```
numerical: 1.866588 analytic: 1.866588, relative error: 2.870578e-08
numerical: -0.777731 analytic: -0.777731, relative error: 4.913355e-08
numerical: 3.048122 analytic: 3.048122, relative error: 4.932950e-09
numerical: 0.378959 analytic: 0.378958, relative error: 1.656307e-07
numerical: -0.976891 analytic: -0.976891, relative error: 3.352302e-08
numerical: -1.707841 analytic: -1.707841, relative error: 1.244064e-08
numerical: -0.748268 analytic: -0.748268, relative error: 2.655352e-08
numerical: 0.381892 analytic: 0.381892, relative error: 1.218675e-07
numerical: 0.044736 analytic: 0.044736, relative error: 9.387103e-08
numerical: 0.417072 analytic: 0.417072, relative error: 3.609779e-09
numerical: 4.627194 analytic: 4.626463, relative error: 7.895356e-05
numerical: -0.706533 analytic: -0.707054, relative error: 3.688338e-04
numerical: 0.229622 analytic: 0.234915, relative error: 1.139498e-02
numerical: 0.643648 analytic: 0.642792, relative error: 6.650322e-04
numerical: -2.861869 analytic: -2.868403, relative error: 1.140199e-03
numerical: 1.028699 analytic: 1.023023, relative error: 2.766471e-03
numerical: 3.339658 analytic: 3.338466, relative error: 1.784581e-04
numerical: -1.248134 analytic: -1.256357, relative error: 3.283001e-03
numerical: -0.407527 analytic: -0.404488, relative error: 3.742294e-03
numerical: -0.759588 analytic: -0.763481, relative error: 2.555840e-03
```

```
In [5]: # Now that we have a naive implementation of the softmax loss function and its gradient,
# implement a vectorized version in softmax_loss_vectorized.
# The two versions should compute the same results, but the vectorized version should
# much faster.
tic = time.time()
loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))

from cs682.classifiers.softmax import softmax_loss_vectorized
tic = time.time()
```

```

loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

# As we did for the SVM, we use the Frobenius norm to compare the two versions
# of the gradient.
grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))
print('Gradient difference: %f' % grad_difference)

```

```

naive loss: 2.355978e+00 computed in 0.140424s
vectorized loss: 2.355978e+00 computed in 0.006468s
Loss difference: 0.000000
Gradient difference: 0.000000

```

```

In [6]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of over 0.35 on the validation set.
from cs682.classifiers import Softmax
results = {}
best_val = -1
best_softmax = None
learning_rates = [1e-7, 5e-7]
regularization_strengths = [2.5e4, 5e4]

softmax1 = Softmax()
for i in learning_rates:
    for j in regularization_strengths:
        loss_history = softmax1.train(X_train, y_train, learning_rate=i, reg=j, num_iter=1000)
        y_train_prediction = softmax1.predict(X_train)
        accuracy_train = np.mean(y_val == y_train_prediction)
        y_val_prediction = softmax1.predict(X_val)
        accuracy_validation = np.mean(y_val == y_val_prediction)

        results[(i,j)] = accuracy_train, accuracy_validation
        if accuracy_validation > best_val:
            best_val=accuracy_validation
            best_softmax = softmax1

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength.
# This should be identical to the validation that you did for the SVM; save
# the best trained softmax classifier in best_softmax.
#####
# Your code
#####

```

```

#                                     END OF YOUR CODE                                     #
#####

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val)

iteration 0 / 1500: loss 761.533042
iteration 100 / 1500: loss 460.528496
iteration 200 / 1500: loss 279.488481
iteration 300 / 1500: loss 169.764818
iteration 400 / 1500: loss 103.437518
iteration 500 / 1500: loss 63.507726
iteration 600 / 1500: loss 39.245693
iteration 700 / 1500: loss 24.554745
iteration 800 / 1500: loss 15.727254
iteration 900 / 1500: loss 10.341743
iteration 1000 / 1500: loss 7.065417
iteration 1100 / 1500: loss 5.082982
iteration 1200 / 1500: loss 3.942867
iteration 1300 / 1500: loss 3.249523
iteration 1400 / 1500: loss 2.783591

/home/akhila/.local/lib/python3.6/site-packages/ipykernel_launcher.py:17: DeprecationWarning: c

iteration 0 / 1500: loss 3.140750
iteration 100 / 1500: loss 2.529964
iteration 200 / 1500: loss 2.324758
iteration 300 / 1500: loss 2.236533
iteration 400 / 1500: loss 2.181742
iteration 500 / 1500: loss 2.148546
iteration 600 / 1500: loss 2.224669
iteration 700 / 1500: loss 2.113984
iteration 800 / 1500: loss 2.085940
iteration 900 / 1500: loss 2.176372
iteration 1000 / 1500: loss 2.196403
iteration 1100 / 1500: loss 2.122962
iteration 1200 / 1500: loss 2.182554
iteration 1300 / 1500: loss 2.224660
iteration 1400 / 1500: loss 2.150173
iteration 0 / 1500: loss 2.113960
iteration 100 / 1500: loss 2.193093

```

```

iteration 200 / 1500: loss 2.140415
iteration 300 / 1500: loss 2.194579
iteration 400 / 1500: loss 2.105710
iteration 500 / 1500: loss 2.055973
iteration 600 / 1500: loss 2.085644
iteration 700 / 1500: loss 2.152067
iteration 800 / 1500: loss 2.146090
iteration 900 / 1500: loss 2.145500
iteration 1000 / 1500: loss 2.063160
iteration 1100 / 1500: loss 2.007901
iteration 1200 / 1500: loss 2.152323
iteration 1300 / 1500: loss 2.196219
iteration 1400 / 1500: loss 2.050577
iteration 0 / 1500: loss 2.344331
iteration 100 / 1500: loss 2.176658
iteration 200 / 1500: loss 2.134563
iteration 300 / 1500: loss 2.111681
iteration 400 / 1500: loss 2.143278
iteration 500 / 1500: loss 2.215188
iteration 600 / 1500: loss 2.184355
iteration 700 / 1500: loss 2.197539
iteration 800 / 1500: loss 2.186359
iteration 900 / 1500: loss 2.190039
iteration 1000 / 1500: loss 2.212629
iteration 1100 / 1500: loss 2.184437
iteration 1200 / 1500: loss 2.244856
iteration 1300 / 1500: loss 2.225584
iteration 1400 / 1500: loss 2.160992
lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.000000 val accuracy: 0.358000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.000000 val accuracy: 0.342000
lr 5.000000e-07 reg 2.500000e+04 train accuracy: 0.000000 val accuracy: 0.356000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.000000 val accuracy: 0.336000
best validation accuracy achieved during cross-validation: 0.358000

```

```

In [7]: # evaluate on test set
        # Evaluate the best softmax on test set
        y_test_pred = best_softmax.predict(X_test)
        test_accuracy = np.mean(y_test == y_test_pred)
        print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))

```

```
softmax on raw pixels final test set accuracy: 0.328000
```

### Inline Question - True or False

It's possible to add a new datapoint to a training set that would leave the SVM loss unchanged, but this is not the case with the Softmax classifier loss.

Your answer: True

*Your explanation:* In SVM, the decision boundary is not significantly effected by one training point. For softmax loss, the exponential of a small value can cause large changes in loss.

```
In [8]: # Visualize the learned weights for each class
w = best_softmax.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)

w_min, w_max = np.min(w), np.max(w)

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])
```

